

Vector similarity search helps us find text that is most similar to a query.

A PDF file usually contains unstructured text, so we first extract the text page by page.

After extraction, we split the text into smaller chunks and convert those chunks into vectors.

Then we compare a user query against those chunk vectors using cosine similarity.